

The Gate of the AOSP #3 : Externals & Extras

## QEMU and Valgrind : Emulator & Memory Technology



2012. 10. 26.

[www.kandroid.org](http://www.kandroid.org) 운영자 : 양정수 (yangjeongsoo at gmail.com), 닉네임: 들풀 🌱

## QEMU and Valgrind : **Emulator & Memory Technology**

### 1. **Background History and Knowledge**

- Android Emulator vs. iPhone Simulator
- Dynamic Binary Translation and Instrumentation
- Memory Issues : Fragmentation and Leakage

### 2. **QEMU : Emulation vs. Performance**

- What is QEMU?
- Goldfish Linux Kernel and Supported ABI(armeabi,armeabi-v7a,x86,mips)
- QEMU Build System : Android Build System, Cygwin, MinGW
- Emulation + Simulation(?) : OpenGL GPU Emulation
- Emulation + H/W Virtualization : x86 and Google TV Emulator
- Some Advanced Topics
  - LLVM Backend for QEMU
  - Can we dream about real s/w phone or cloud phone ?

### 3. **Valgrind : Memory Leak vs. Memory Management**

- What is Valgrind?
- Valgrind Build and Installation for Real Target
- Valgrind Issues : No Swap and Low Memory Killer
- Memory Analysis Tools & Memory Management Methods  
(Java:MAT, JNI:Global/Local Ref, C++:sp,wp, SysV Shmem)
- Some Advanced Topics
  - Valgrind and LLVM's AddressSanitizer
  - What Every Android Programmer Should Know About Memory?


# Android Emulator vs. iPhone Simulator

## Emulation versus simulation

The word "emulator" was coined in 1963 at IBM<sup>[15]</sup> during development of the NPL (IBM 360) product line, using a "new combination of software, microcode, and hardware".<sup>[16]</sup> They discovered that using [microcode](#) hardware instead of software simulation, to execute programs written for earlier IBM computers, dramatically increased simulation speed. Earlier in 1957, IBM provided the IBM 709 computer with an [interpreter](#) program (software) to execute legacy programs written for the [IBM 704](#) to run on the [IBM 709](#) and later on the IBM 7090<sup>[17]</sup> In 1963, when microcode was first used to speed up this simulation process, IBM engineers coined the term "emulator" to describe the concept.

It has recently become common to use the word "emulate" in the context of software. However, before 1980, "emulation" referred only to emulation with a hardware or microcode assist, while "simulation" referred to pure software emulation.<sup>[18]</sup> For example, a computer specially built for running programs designed for another architecture is an emulator. In contrast, a simulator could be a program which runs on a PC, so that old Atari games can be simulated on it. Purists continue to insist on this distinction, but currently the term "emulation" often means the complete imitation of a machine executing binary code while "simulation" often refers to [Computer simulation](#), where a computer program is used to simulate an abstract model. Computer simulation is used in virtually every scientific and engineering domain and Computer Science is no exception, with several projects simulating abstract models of computer systems, such as [Network simulation](#).

source : [http://en.wikipedia.org/wiki/Emulator#Emulation\\_vs\\_simulation](http://en.wikipedia.org/wiki/Emulator#Emulation_vs_simulation)



Why dynamic, not static ?

- Static method is very difficult (ex, indirect branches)

Dynamic Binary Translation :

- ex) Translating arm executable into x86 executable


Dynamic Binary Instrumentation :

- Injecting arbitrary code into binary at runtime.

| Dynamic Binary Translation   | Dynamic Binary Instrumentation   |
|--|--|
| <div>Simple Translation</div> <ul style="list-style-type: none"><li>• fetch-decode-execute loop</li></ul> <div>Advanced Translation</div> <ul style="list-style-type: none"><li>• Caching</li><li>• Recompilation (ex, JIT Compiler)</li></ul> | <div>Decompilation &amp; Recompilation</div> <ul style="list-style-type: none"><li>• Valgrind</li></ul> <div>Using Trap Instruction</div> <ul style="list-style-type: none"><li>• Dtrace</li><li>• SystemTap</li><li>• Frysk</li><li>• GDB</li></ul> |

# Dynamic Binary Translation and Instrumentation : fetch-decode-execute loop

---



0xc0083420: e0861005

```
add r1, r6, r5
```

```
switch(opcode) {
```

```
case OP_ADD:
```

```
    add(r1, r6, r5);
```

```
    break;
```

```
}
```

```
void add(int r1, int r6, int r5)
```

```
{
```

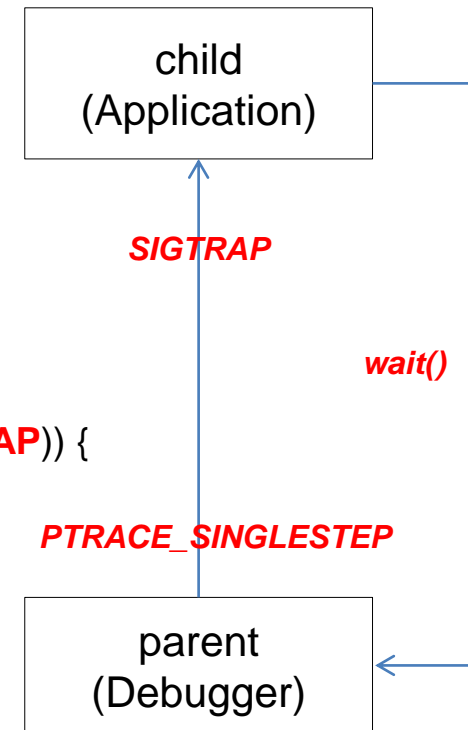
```
    context->reg[r1] = context->reg[r6]  
                      + context->reg[r5];
```

```
}
```

# Dynamic Binary Translation and Instrumentation : trap instruction

```
int main(int argc, char **argv)
{
    pid_t child;

    child = fork();
    if(child == 0) {
        ptrace(PTRACE_TRACEME, 0, NULL, NULL);
        execl(argv[1], argv[1], NULL);
    }
    else {
        int status;
        ptrace(PTRACE_SINGLESTEP, child, NULL, NULL);
        while(1) {
            wait(&status);
            if (WIFSTOPPED(status) && (WSTOPSIG(status) == SIGTRAP)) {
                ptrace(PTRACE_SINGLESTEP, child, NULL, NULL);
            }
            else {
                break;
            }
        }
    }
    return 0;
}
```



# Memory Issues : Fragmentation and Leakage

## Memory Fragmentation

### Internal Fragmentation

- Linux kernel : slab allocator
- Libc : dlmalloc, ptmalloc

### External Fragmentation

- Linux kernel : buddy system

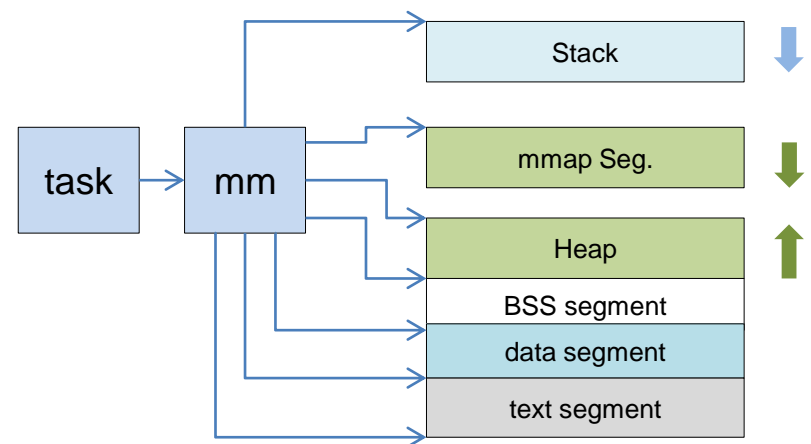
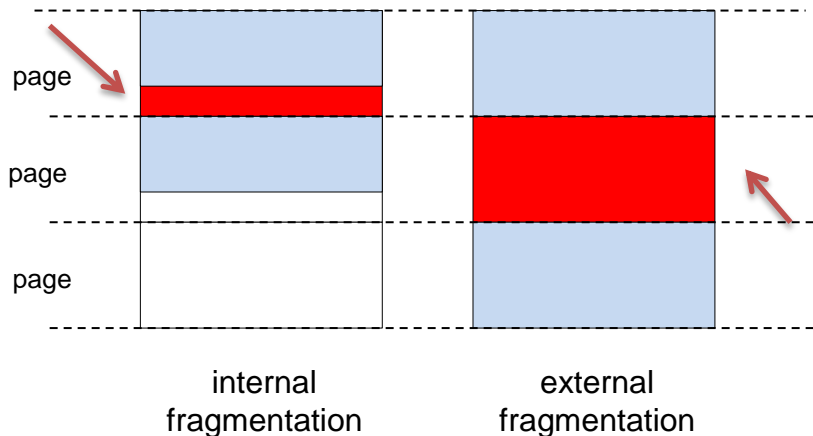
## Memory Leakage

### Memory Allocation

- Static allocation
- Stack allocation
- Heap(Dynamic) allocation

### Garbage Collection

Explicit collection  
Automatic collection



## QEMU and Valgrind : **Emulator & Memory Technology**

### 1. **Background History and Knowledge**

- Android Emulator vs. iPhone Simulator
- Dynamic Binary Translation and Instrumentation
- Memory Issues : Fragmentation and Leakage

### 2. **QEMU : Emulation vs. Performance**

- What is QEMU?
- Goldfish Linux Kernel and Supported ABI(armeabi,armeabi-v7a,x86,mips)
- QEMU Build System : Android Build System, Cygwin, MinGW
- Emulation + Simulation(?) : OpenGL GPU Emulation
- Emulation + H/W Virtualization : x86 and Google TV Emulator
- Advanced Topics
  - LLVM Backend for QEMU
  - Can we dream about real s/w phone or cloud phone ?

### 3. **Valgrind : Memory Leak vs. Memory Management**


- What is Valgrind?
- Valgrind Build and Installation for Real Target
- Valgrind Issues : No Swap and Low Memory Killer
- Memory Analysis Tools & Memory Management Methods (Java:MAT, JNI:Global/Local Ref, C++:sp,wp, SysV Shmem)
- Advanced Topics
  - Valgrind and LLVM's AddressSanitizer
  - What Every Android Programmer Should Know About Memory?



# What is QEMU ?

---

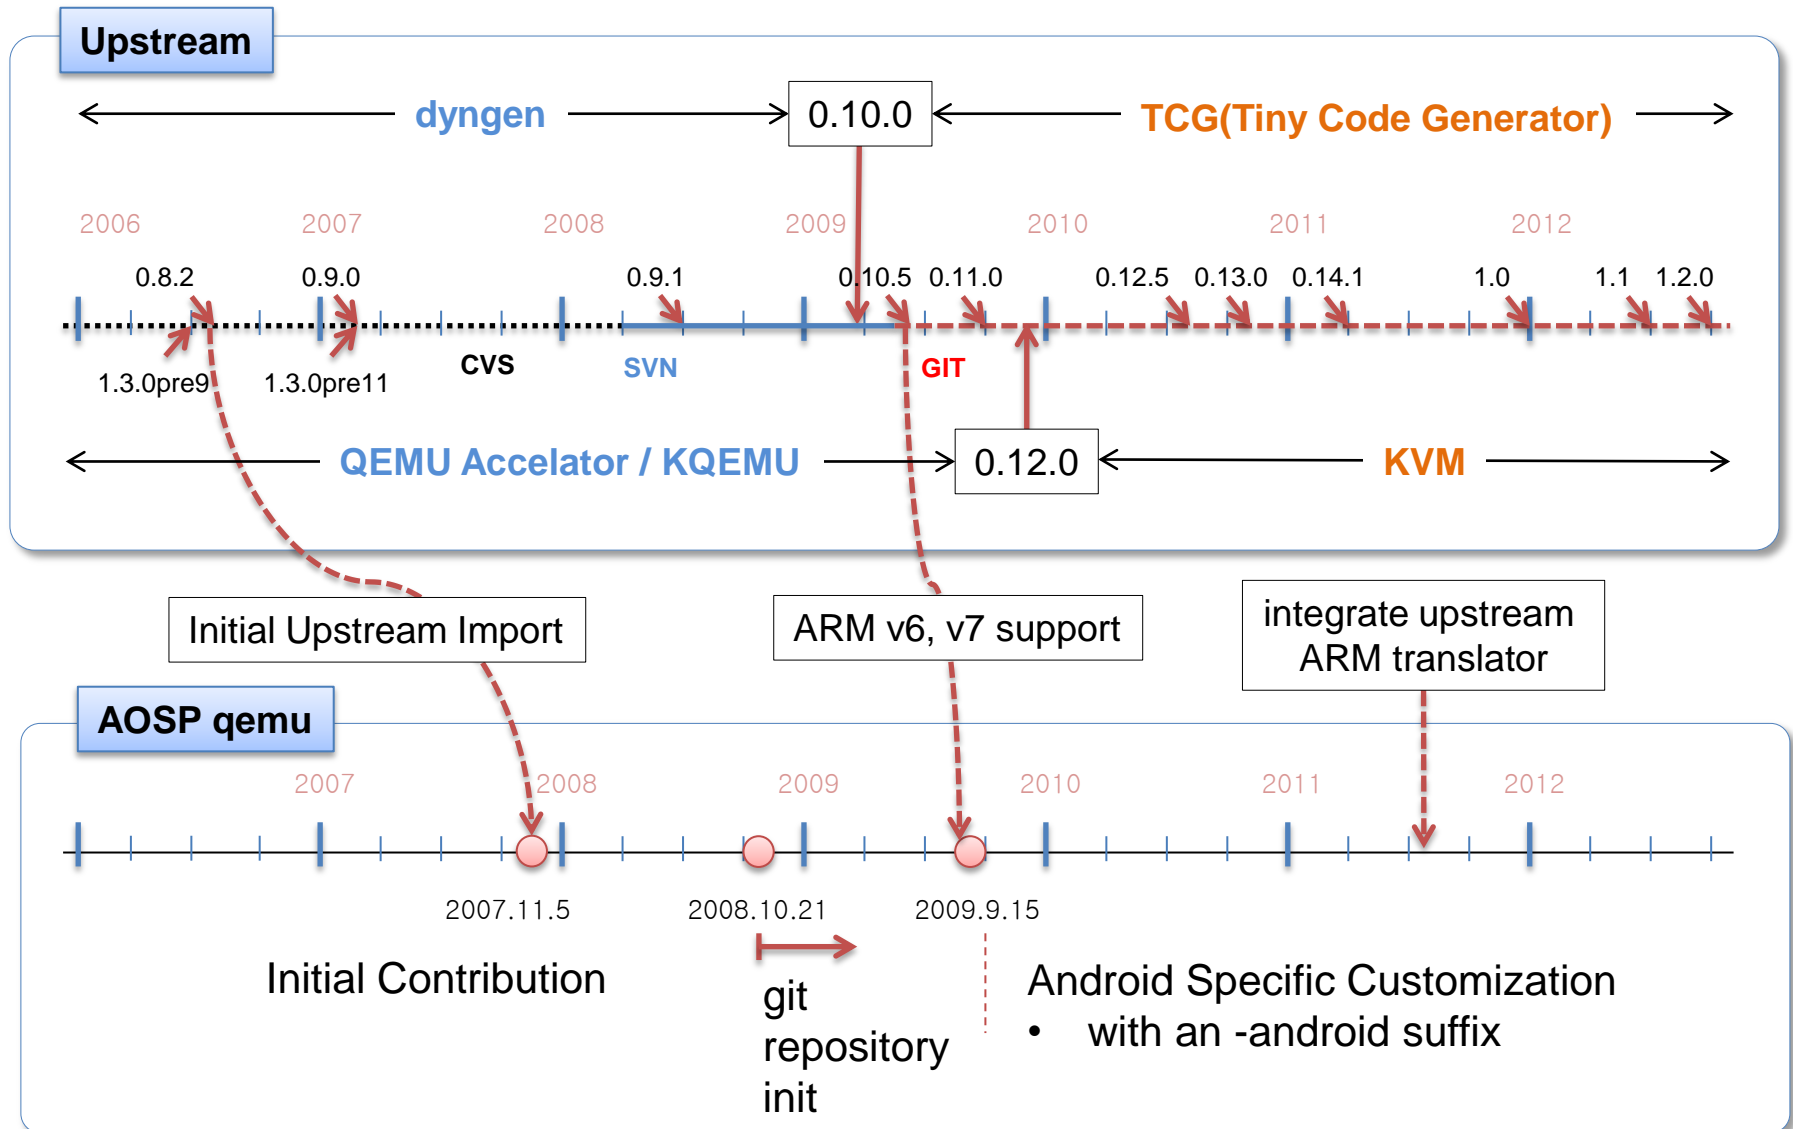
- Open source library for hardware **emulation** and **virtualization**
- Fast CPU and device emulator based on **dynamic binary translation**
- Execution of SW binaries of a guest instruction set on host PC
  - Development of drivers and application SW on host PC
  - Debugging of guest binaries
  - Development of SW tool chain for guest SW
- Different operation modes
  - **full system emulation**
  - user mode emulation
  - **KQEMU/KVM**



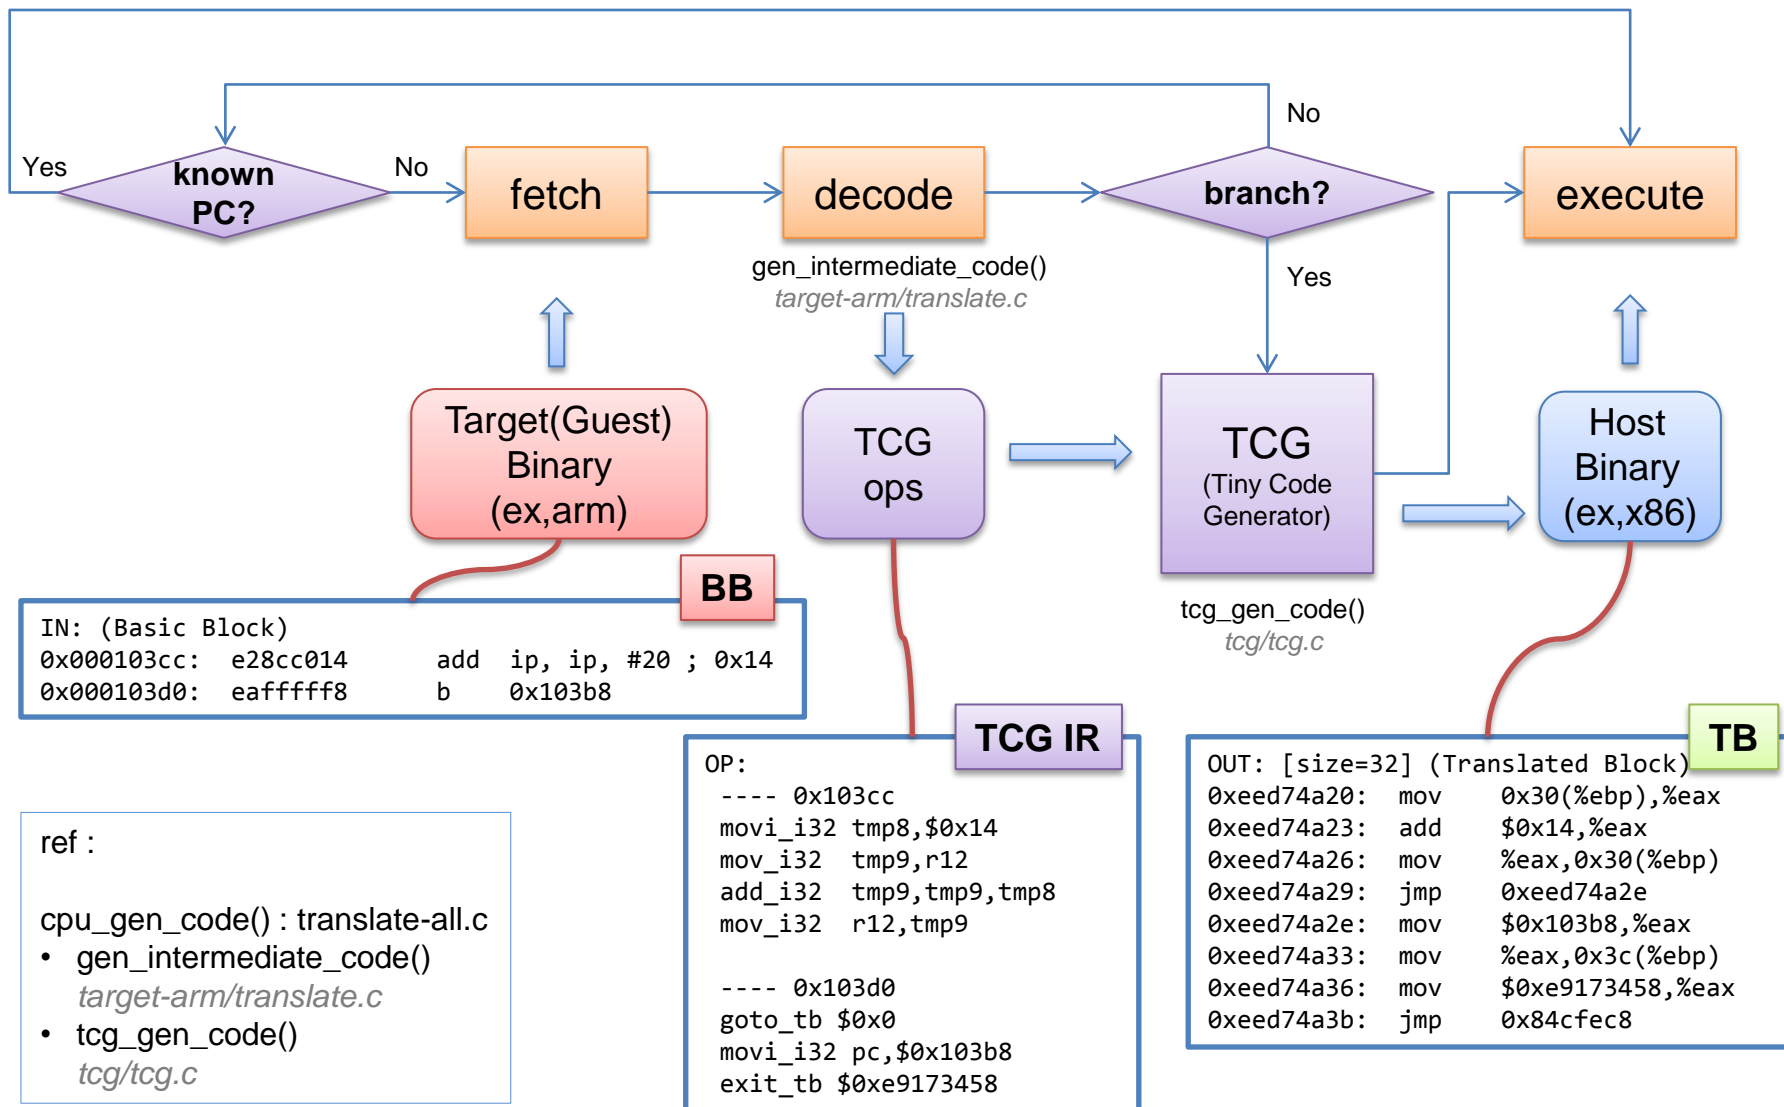
**Fabrice Bellard**  
born :1972  
Grenoble, France

Source : *First QEMU Users' Forum* - 1st International QEMU Users' Forum, 2011

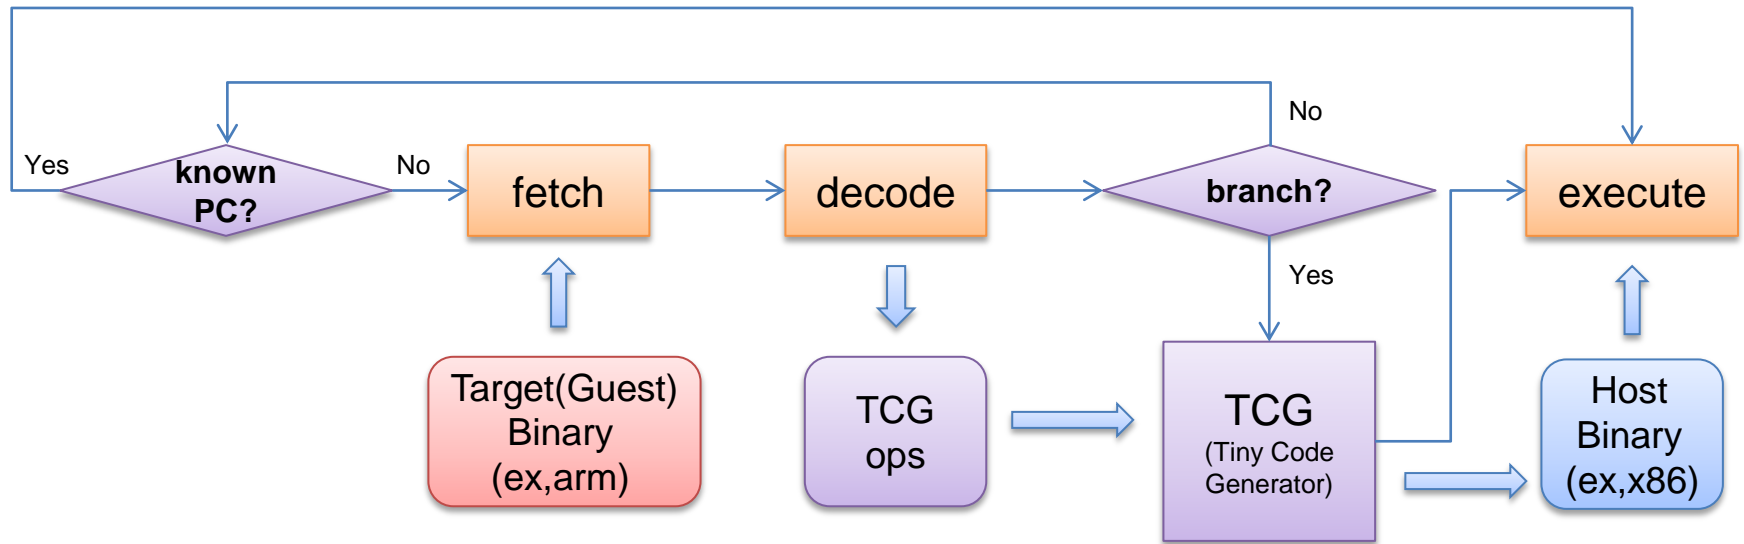
# What is QEMU ?



# What is QEMU ?



# What is QEMU ?

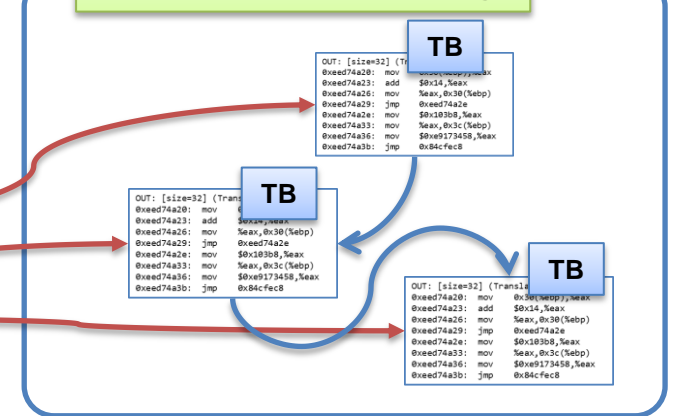


```
for (;;) {
    tb = tb_find_fast();
}
```

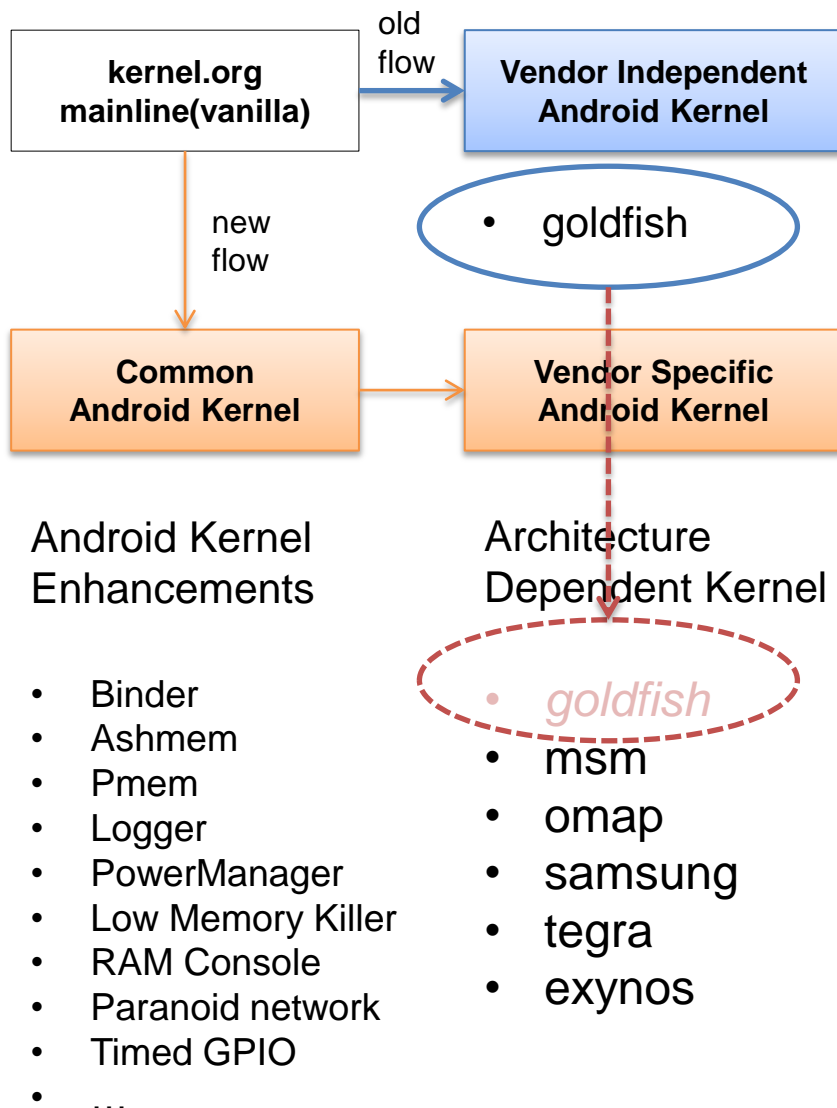
Hash Table  
(tb\_jmp\_cache)

|        |
|--------|
| 0xa7a1 |
| 0xd37c |
| 0x4389 |
| ...    |

Translation Cache  
(Direct Block Chaining)



# Goldfish Linux Kernel and Supported ABI(armeabi,armeabi-v7a,x86,mips)



For the first android SDK,

1. Initial qemu Upstream Import
2. Initial contribution

What is the initial contribution?

1. Android-specific modification to qemu
2. Create virtual soc : goldfish
  - Support CPU : ARM926EJS
  - Support Device : fb,audio,tty,mmc,nand, etc

What is the add-on features?

1. Support additional ABI : armv7, x86
2. OpenGL ES h/w acceleration
3. Emulator with h/w assist for x86

## Android Build System

- AOSP Build System : GNU Make + Toolchain + Custom Build Tools
- NDK Build System : **Cygwin** + Toolchain + Custom Build Tools

## How can we build android emulator?

```
$ cd ~/asop/external/qemu; ./android-configure.sh
```

- for Linux

```
$ ./android-rebuild.sh
```
- for Windows

```
$ ./android-rebuild.sh --static --mingw
```

## Cygwin vs. MinGW

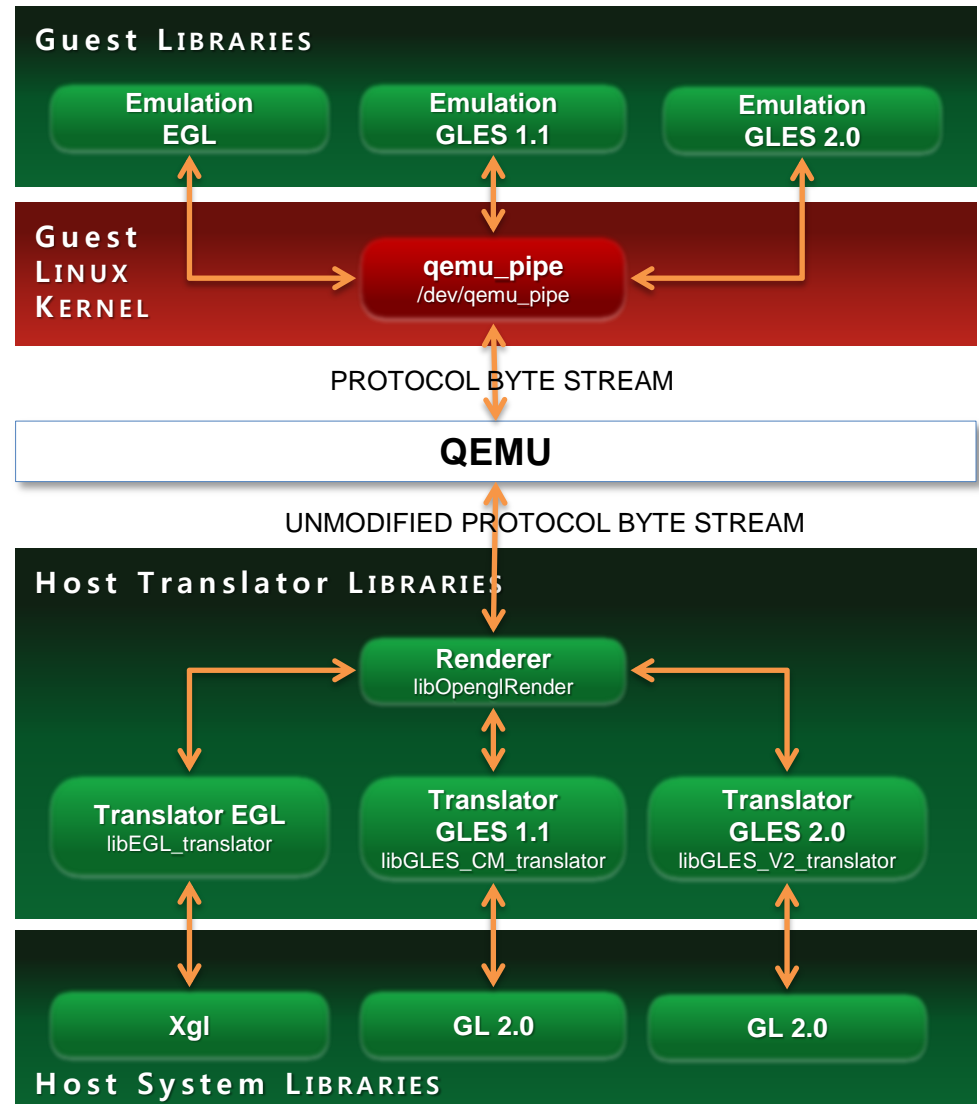
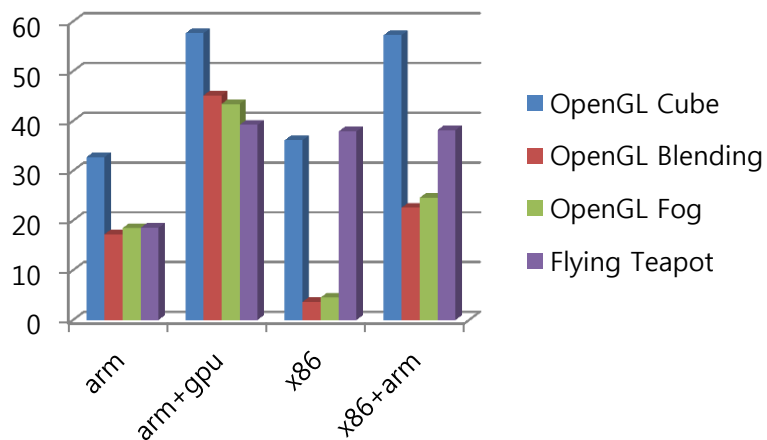
- The **Cygwin tools** are ports of the popular GNU development tools for Microsoft Windows. They run thanks to the Cygwin library which provides the POSIX system calls and environment these programs expect.
- **MinGW** ("Minimalistic GNU for Windows") is a collection of freely available and freely distributable Windows specific header files and import libraries combined with GNU toolsets that allow one to produce native Windows programs **that do not rely on any 3rd-party C runtime DLLs.**

# Emulation + Simulation(?) : OpenGL GPU Emulation



|                 | arm       | arm+gpu  | x86      | x86+arm  |
|-----------------|-----------|----------|----------|----------|
| OpenGL Cube     | 32.726646 | 57.63191 | 36.11239 | 57.23642 |
| OpenGL Blending | 17.244913 | 45.1264  | 3.632293 | 22.54601 |
| OpenGL Fog      | 18.457346 | 43.40562 | 4.50357  | 24.5354  |
| Flying Teapot   | 18.511148 | 39.20031 | 37.88435 | 38.16494 |

<http://code.google.com/p/0xbench/>



# Emulation + H/W Virtualization : x86 and Google TV Emulator

## Google TV Emulator install process

### 1. Check System Requirement : Linux OS with KVM

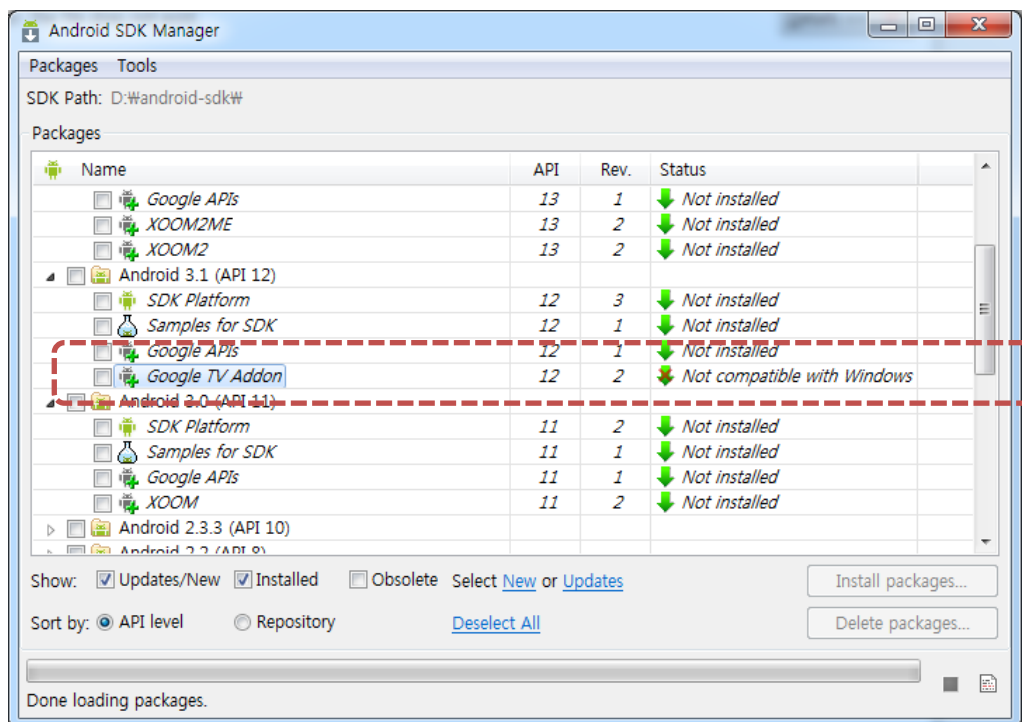
```
$ cat /proc/cpuinfo
$ egrep '^flags.*(vmx|svm)' /proc/cpuinfo
$ sudo modprobe kvm-intel (or sudo modprobe kvm-amd)
```



### 2. Linux KVM Install Process

```
$ make
$ make modules
$ make modules_install

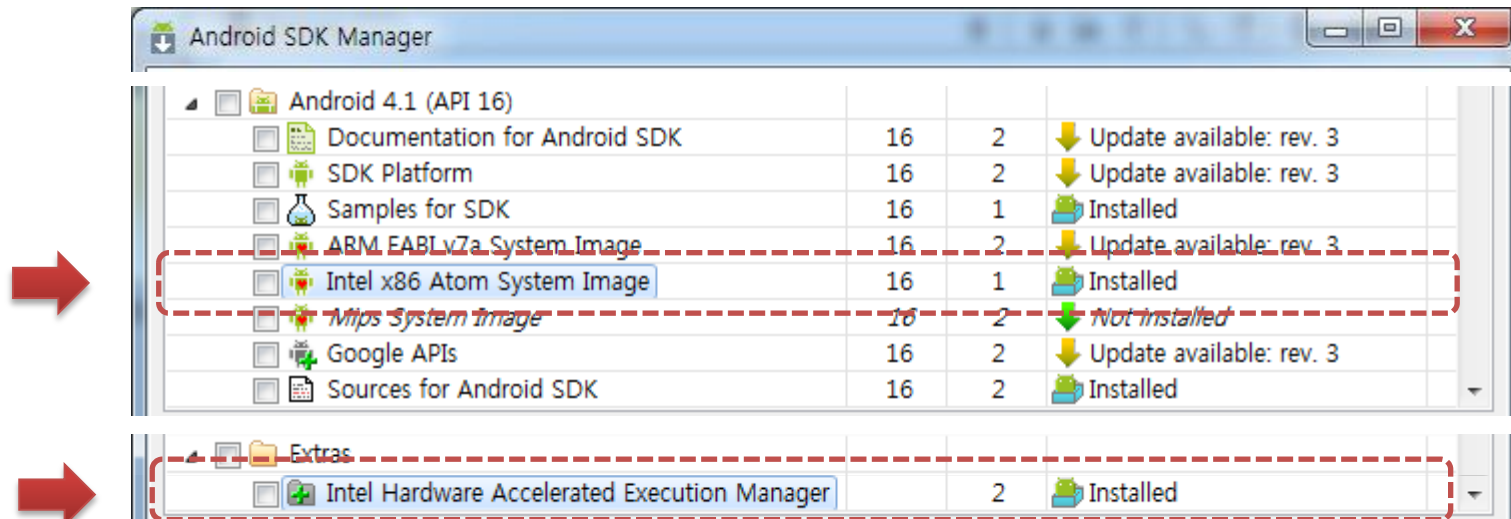
$ sudo modprobe kvm-intel
$ ls -l /dev/kvm
$ sudo chmod a+rw /dev/kvm
$ sudo addgroup kvmusers
$ sudo addgroup you kvmusers
$ sudo chgrp kvmusers /dev/kvm
$ sudo chmod g+rw /dev/kvm
$ ls -l /dev/kvm
```





# Emulation + H/W Virtualization : x86 and Google TV Emulator

## Intel HAX based Android Emulator Acceleration



```
D:\android-sdk\tools>sc query intelhaxm

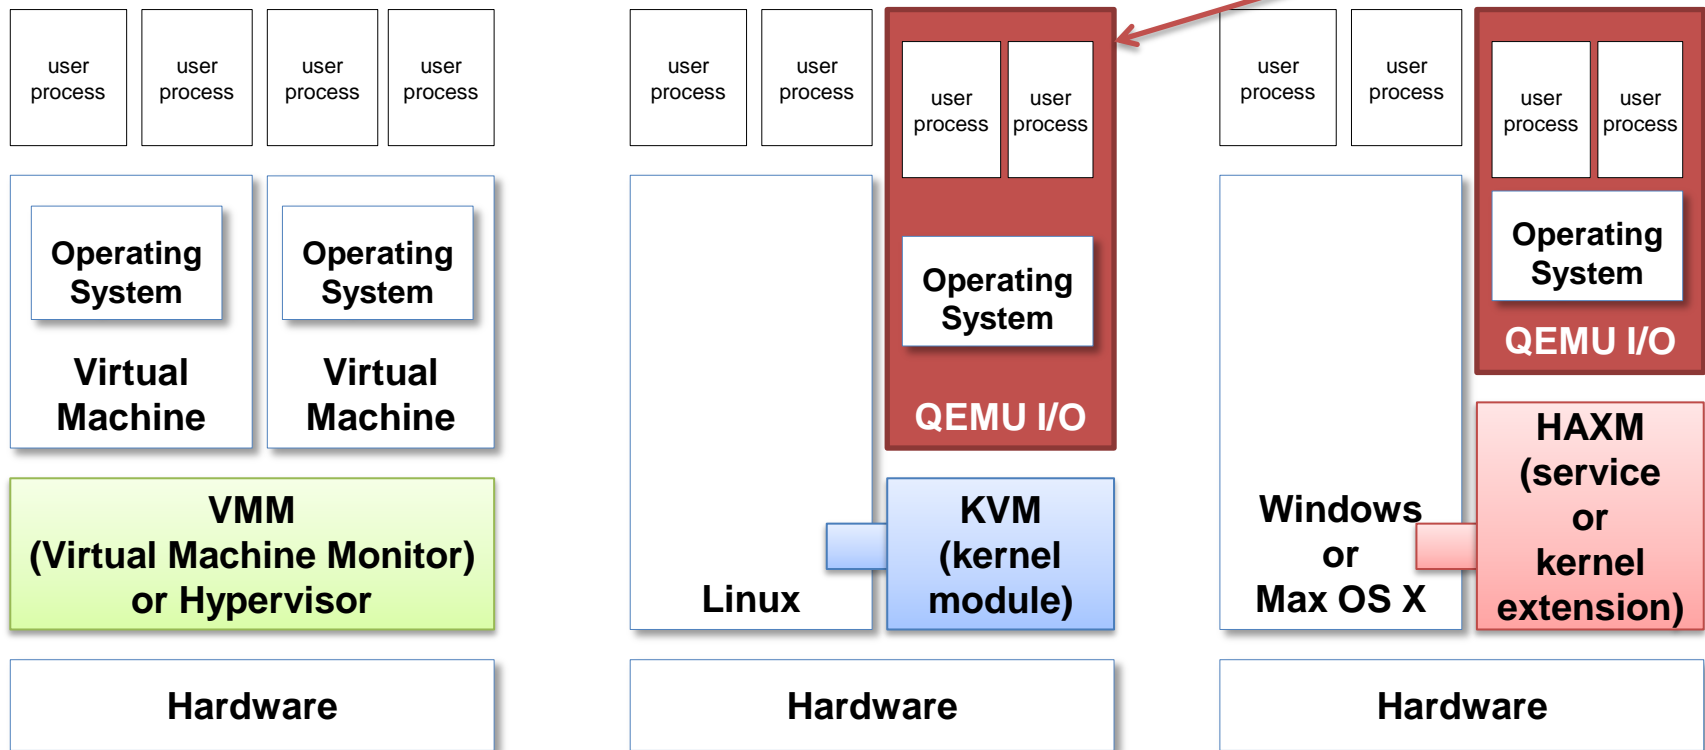
SERVICE_NAME: intelhaxm
        종 류               : 1  KERNEL_DRIVER
        상 태               : 4  RUNNING
                           (STOPPABLE, NOT_PAUSABLE, IGNORES_SHUTDOWN)
        WIN32_EXIT_CODE      : 0  (0x0)
        SERVICE_EXIT_CODE   : 0  (0x0)
        검사점              : 0x0
        WAIT_HINT            : 0x0

D:\android-sdk\tools>
```

# Emulation + H/W Virtualization : x86 and Google TV Emulator

What is KVM, Intel HAXM?

- KVM is a Hypervisor/VMM based on Linux
- HAXM is a Hypervisor/VMM based on Windows or Max OS X
- KVM adds a third mode to the standard linux **kernel** and **user** modes, the **guest** mode



### LLVM Backend for QEMU

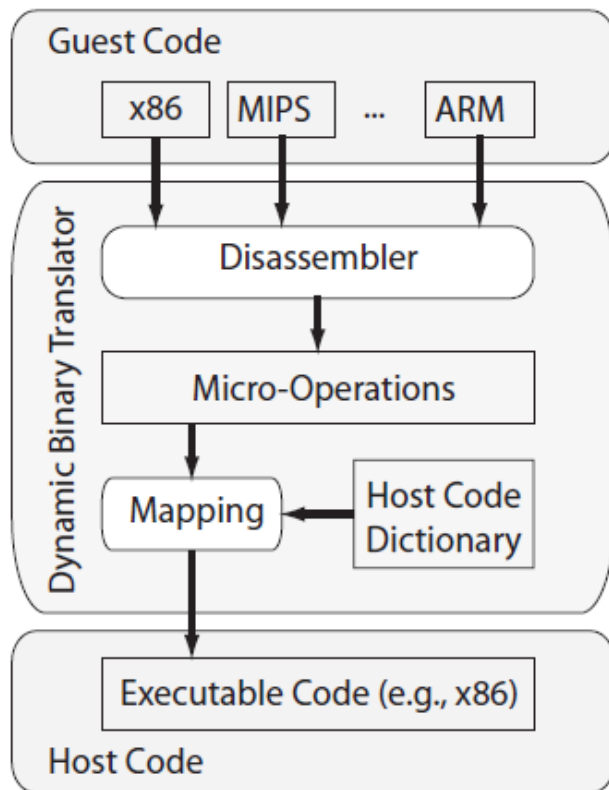


Figure 1: The QEMU dynamic binary translator

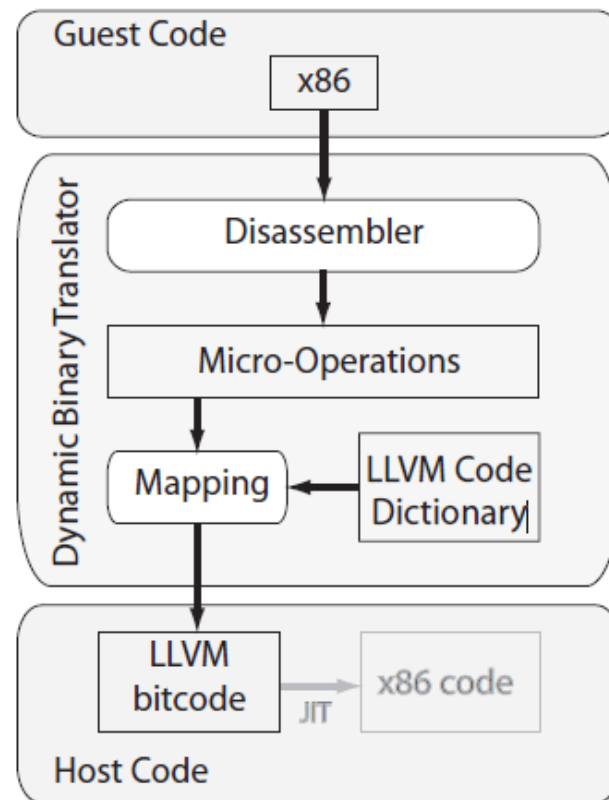


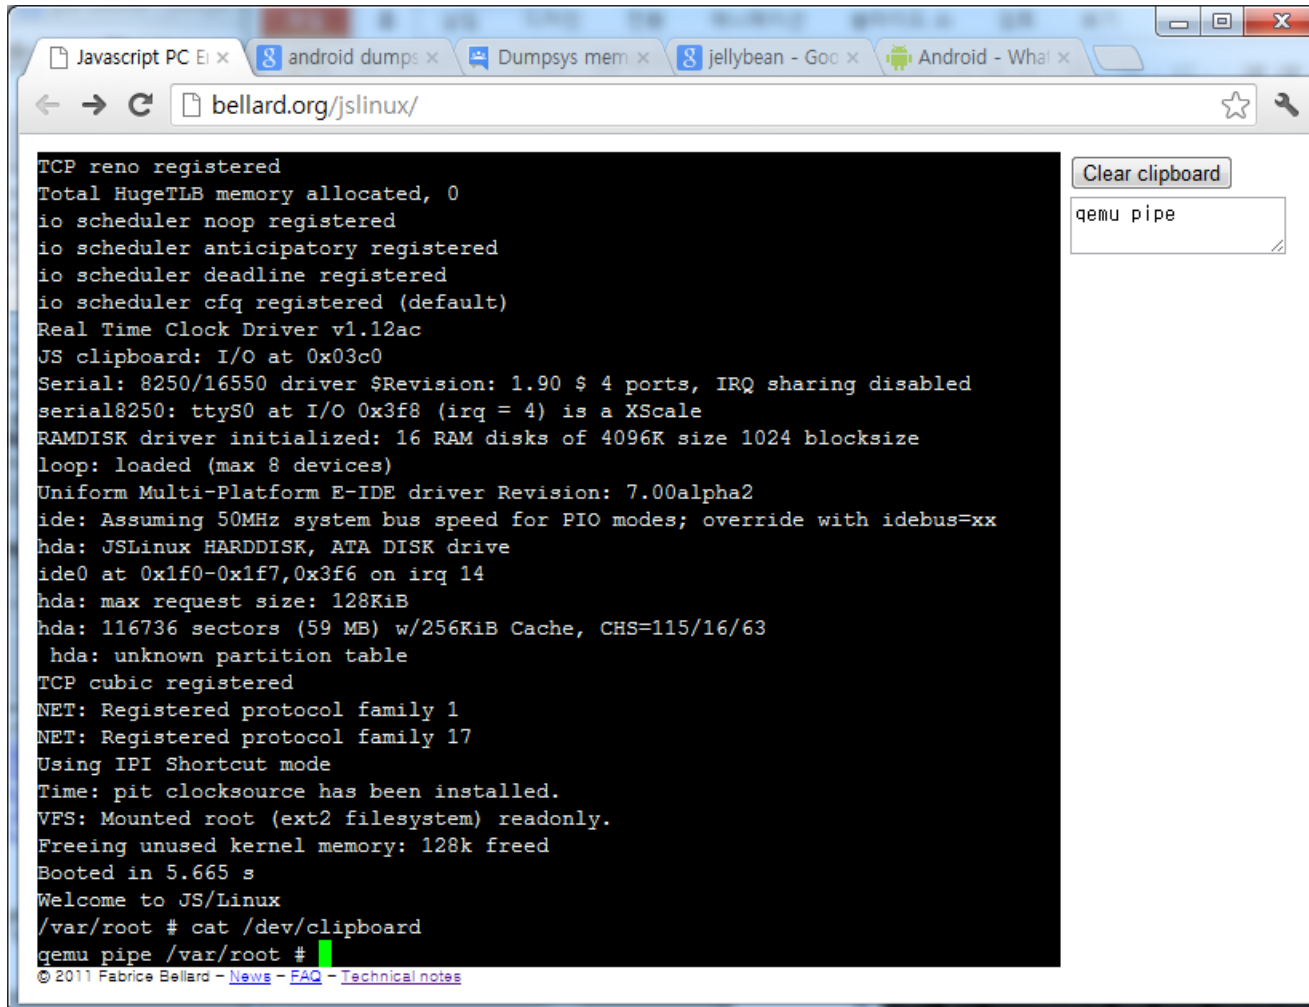
Figure 2: An LLVM backend for QEMU

source : [http://infoscience.epfl.ch/record/149975/files/x86-llvm-translator-chipounov\\_2.pdf](http://infoscience.epfl.ch/record/149975/files/x86-llvm-translator-chipounov_2.pdf)

## Advanced Topics (2)

Can we dream about real s/w phone or cloud phone?

source : <http://bellard.org/jslinux/tech.html>

A screenshot of a web browser window displaying the JS/Linux project page. The browser has several tabs open: 'Javascript PC E...', 'android dumps x', 'Dumpsys mem x', 'jellybean - Goo x', and 'Android - What x'. The address bar shows 'bellard.org/jslinux/'. The main content area is a black terminal window with white text showing the boot process of JS/Linux. The text includes various system initialization messages like 'TCP reno registered', 'Total HugeTLB memory allocated, 0', 'io scheduler noop registered', 'Real Time Clock Driver v1.12ac', 'JS clipboard: I/O at 0x03c0', 'Serial: 8250/16550 driver \$Revision: 1.90 \$ 4 ports, IRQ sharing disabled', 'serial8250: ttyS0 at I/O 0x3f8 (irq = 4) is a XScale', 'RAMDISK driver initialized: 16 RAM disks of 4096K size 1024 blocksize', 'loop: loaded (max 8 devices)', 'Uniform Multi-Platform E-IDE driver Revision: 7.00alpha2', 'ide: Assuming 50MHz system bus speed for PIO modes; override with idebus=xx', 'hda: JSLinux HARDDISK, ATA DISK drive', 'ide0 at 0x1f0-0x1f7,0x3f6 on irq 14', 'hda: max request size: 128KiB', 'hda: 116736 sectors (59 MB) w/256KiB Cache, CHS=115/16/63', 'hda: unknown partition table', 'TCP cubic registered', 'NET: Registered protocol family 1', 'NET: Registered protocol family 17', 'Using IPI Shortcut mode', 'Time: pit clocksource has been installed.', 'VFS: Mounted root (ext2 filesystem) readonly.', 'Freeing unused kernel memory: 128k freed', 'Booted in 5.665 s', 'Welcome to JS/Linux', and a prompt '/var/root # cat /dev/clipboard' followed by 'qemu pipe'. At the bottom of the terminal, it says '© 2011 Fabrice Bellard - News - FAQ - Technical notes'. To the right of the terminal, there is a 'Clear clipboard' button and a text input field containing 'qemu pipe'.

## QEMU and Valgrind : **Emulator & Memory Technology**

### 1. **Background History and Knowledge**

- Android Emulator vs. iPhone Simulator
- Dynamic Binary Translation and Instrumentation
- Memory Issues : Fragmentation and Leakage

### 2. **QEMU : Emulation vs. Performance**

- What is QEMU?
- Goldfish Linux Kernel and Supported ABI(armeabi,armeabi-v7a,x86,mips)
- QEMU Build System : Android Build System, Cygwin, MinGW
- Emulation + Simulation(?) : OpenGL GPU Emulation
- Emulation + H/W Virtualization : x86 and Google TV Emulator
- Advanced Topics
  - LLVM Backend for QEMU
  - Can we dream about real s/w phone or cloud phone ?

### 3. **Valgrind : Memory Leak vs. Memory Management**

- What is Valgrind?
- Valgrind Build and Installation for Real Target
- Valgrind Issues : No Swap and Low Memory Killer
- Memory Analysis Tools & Memory Management Methods  
(Java:MAT, JNI:Global/Local Ref, C++:sp,wp, SysV Shmem)
- Advanced Topics
  - Valgrind and LLVM's AddressSanitizer
  - What Every Android Programmer Should Know About Memory?

# What is Valgrind?

pronounce : “**Val-grinned**”, not “Val-grined”

- Valgrind is a “program-execution monitoring framework”.
- Valgrind comes with many tools, the tool you will use most often is the memcheck tool.
- Memcheck will detect and report the following types of memory errors:
  - Use of uninitialized memory.
  - Reading/writing to memory after it has been freed.
  - Reading/writing off the end of malloc'd blocks.
  - Reading/writing inappropriate areas on the stack.
  - Overlapping src and dest pointers in memcpy() and related functions.
  - other stuff...



Julian Seward is a recycled compiler hacker. he founded the Valgrind project in 2000 and is the project lead and a full time developer. He's also the author of bzip2, a data compression program.



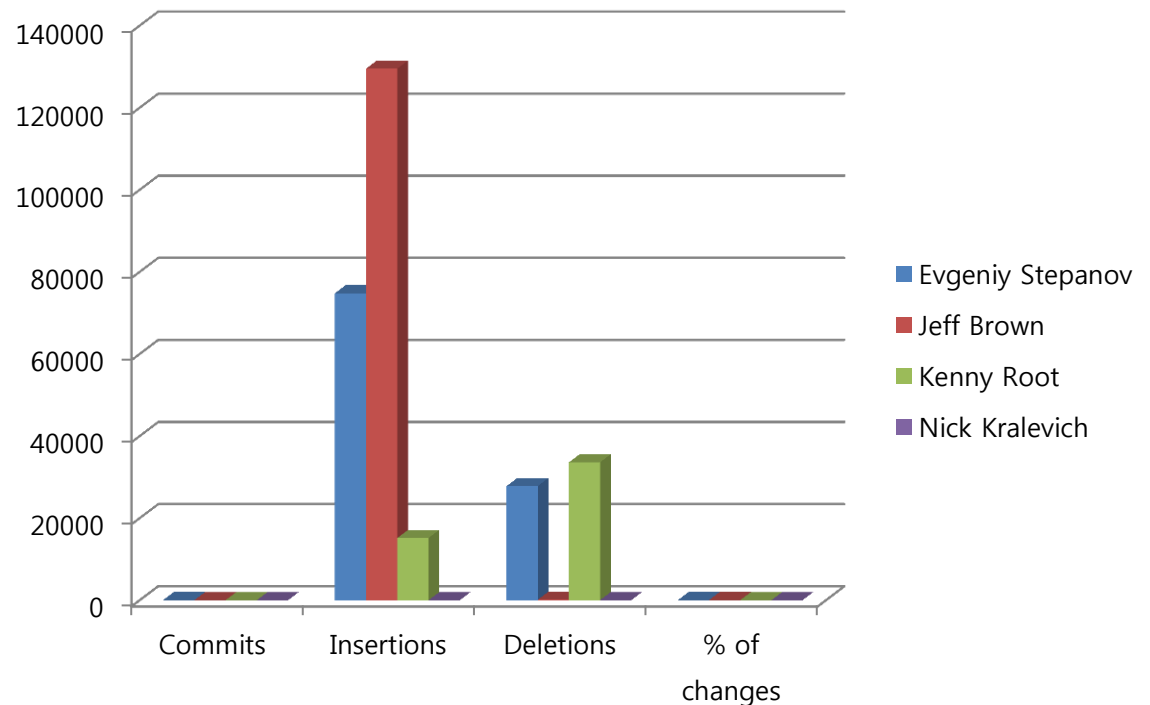
Valgrind, in Norse mythology, is the sacred gate to Valhalla through which only the chosen slain can pass.

Source : <http://ugweb.cs.ualberta.ca/~c201/F06/resources/valgrind.pdf>

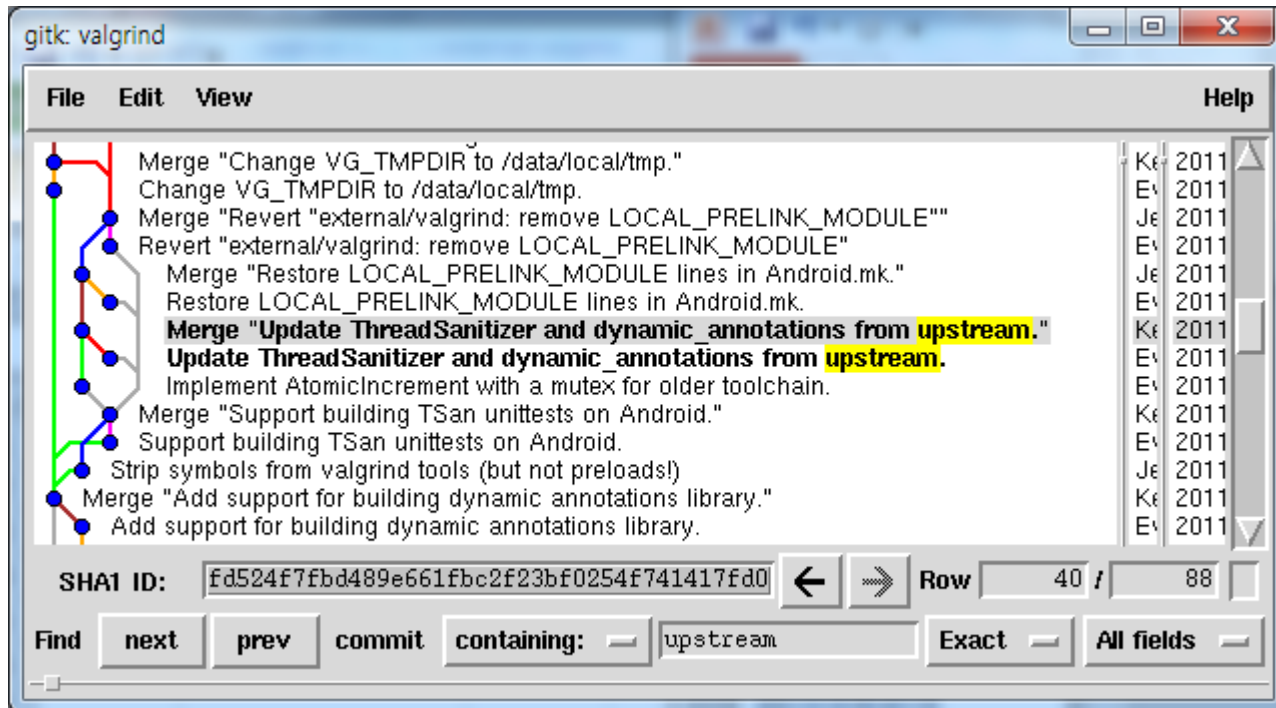
# What is Valgrind?

Initial upstream  
import : valgrind 3.6.0  
- 2011.2.4

| Author                  | Commits | Insertions | Deletions | % of changes |
|-------------------------|---------|------------|-----------|--------------|
| <b>Evgeniy Stepanov</b> | 24      | 74792      | 27816     | 36.51        |
| Jeff Brown              | 4       | 129524     | 123       | 46.13        |
| Kenny Root              | 12      | 15163      | 33595     | 17.35        |
| Nick Kralevich          | 2       | 8          | 14        | 0.01         |



# What is Valgrind?



## Finding races and memory errors with LLVM instrumentation

AddressSanitizer, ThreadSanitizer

Timur Iskhodzhanov, Alexander Potapenko, Alexey Samsonov,  
**Kostya Serebryany**, Evgeniy Stepanov, Dmitriy Vyukov

LLVM Dev Meeting  
November 18, 2011

source : [http://llvm.org/devmtg/2011-11/Serebryany\\_FindingRacesMemoryErrors.pdf](http://llvm.org/devmtg/2011-11/Serebryany_FindingRacesMemoryErrors.pdf)



# Valgrind Build and Installation for Real Target

```
$ cd ~/android-4.1.1_r1/external/valgrind/
```

Modify VG\_LIBDIR cflag.

```
$ vi ~/android-4.1.1_r1/external/valgrind/main/Android.mk
```

```
- -DVG_LIBDIR="/system/lib/valgrind"
```

```
+ -DVG_LIBDIR="/data/local/inst/lib/valgrind" )
```

```
$ mm -B
```

```
-- bin
| |-- racecheck_unittest
| `-- valgrind
|-- lib
|   |-- valgrind
|     |-- cachegrind-arm-linux
|     |-- callgrind-arm-linux
|     |-- default.supp
|     |-- drd-arm-linux
|     |-- helgrind-arm-linux
|     |-- massif-arm-linux
|     |-- memcheck-arm-linux
|     |-- none-arm-linux
|     |-- tsan-arm-linux
|     |-- vgpreload_core-arm-linux.so
|     |-- vgpreload_drd-arm-linux.so
|     |-- vgpreload_helgrind-arm-linux.so
|     |-- vgpreload_massif-arm-linux.so
|     |-- vgpreload_memcheck-arm-linux.so
|     `-- vgpreload_tsan-arm-linux.so
```

```
{
  bionic_dlopen
  Memcheck:Addr4
  fun:_dl_strlen
  ...
  fun:_dl_find_library
  fun:_dl_dlopen
}

{
  bionic_dlopen_c
  Memcheck:Cond
  fun:_dl_strlen
  ...
  fun:_dl_find_library
  fun:_dl_dlopen
}

{
  sha1_block_data_order-reads-below-sp
  Memcheck:Addr4
  fun:sha1_block_data_order
}

{
  dvmPlatformInvoke-misinterpretation-1
  Memcheck:Addr4
  fun:dvmCallJNIMethod_virtualNoRef
}

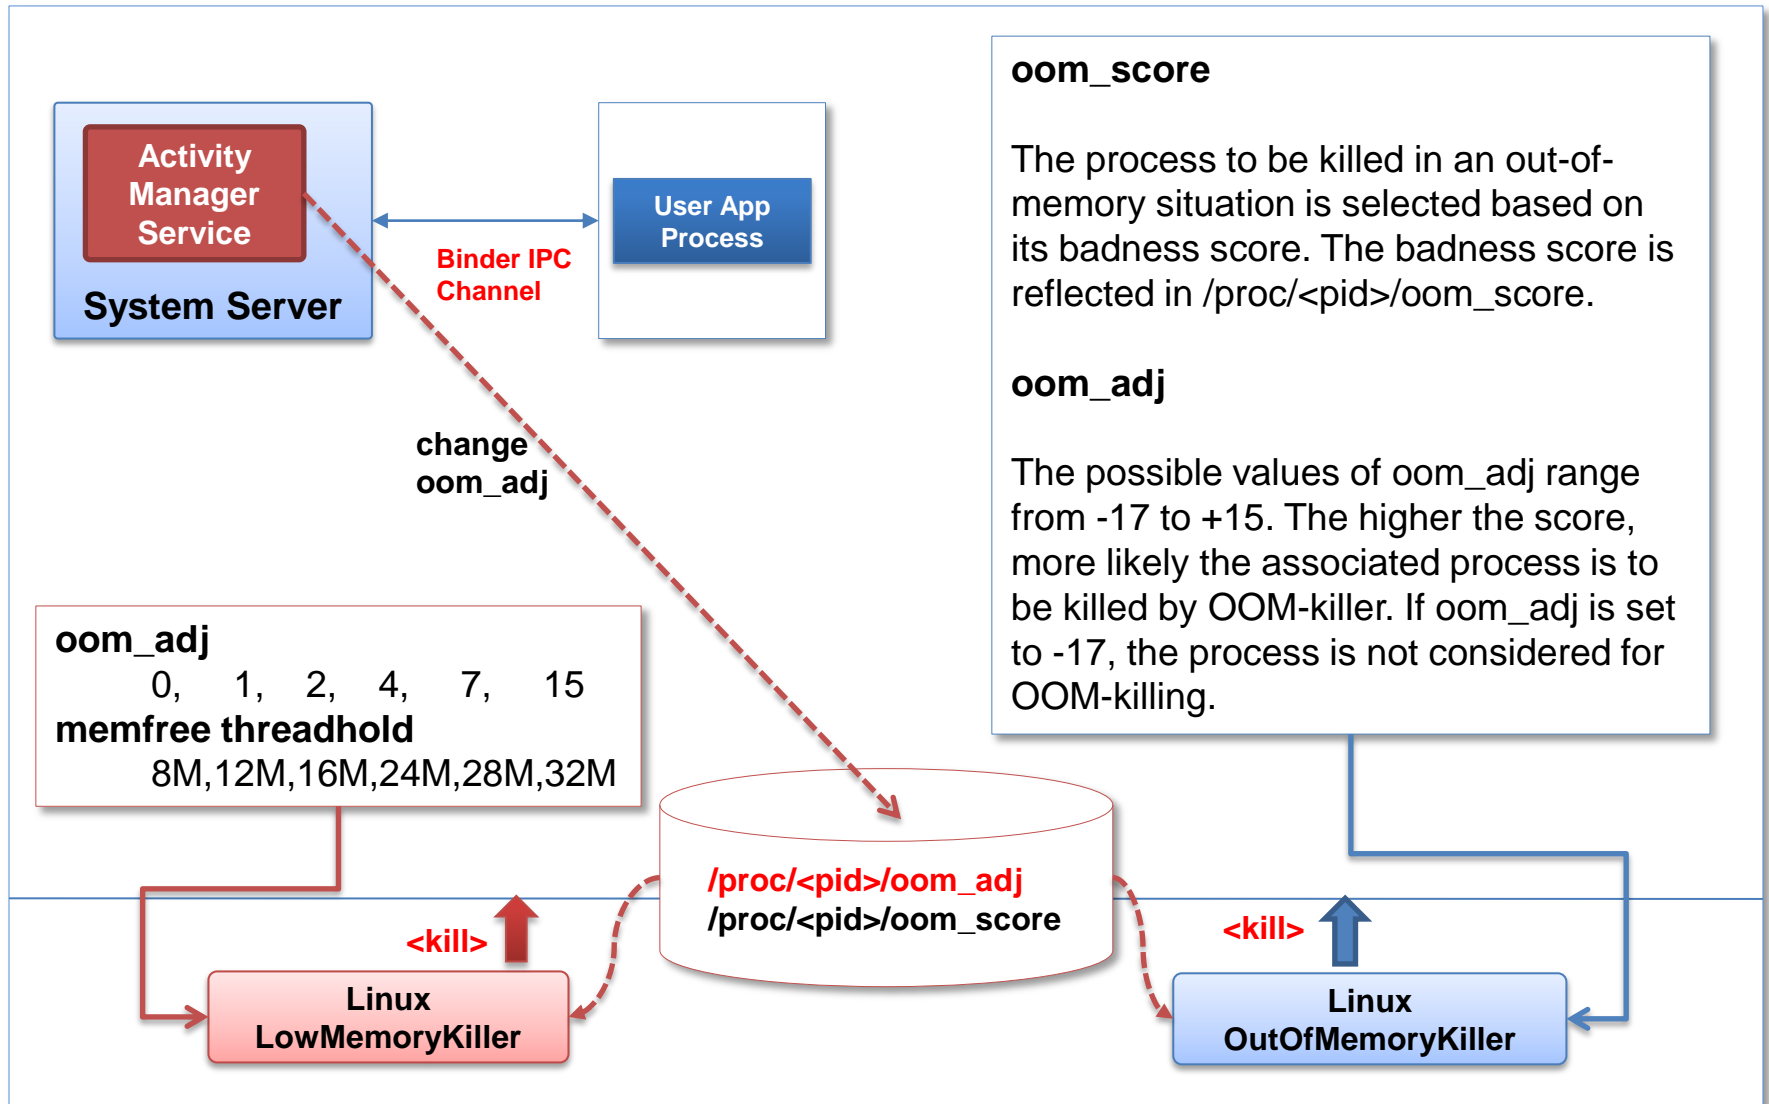
{
  dvmPlatformInvoke-misinterpretation-2
  Memcheck:Addr4
  fun:dvmCallJNIMethod_staticNoRef
}
```

# Valgrind Build and Installation for Real Target



```
C:\Windows\system32\cmd.exe - adb shell
shell@android:/data/local/inst/bin $ getprop ro.product.device
getprop ro.product.device
maguro
shell@android:/data/local/inst/bin $ ./valgrind ls
./valgrind ls
==7345== Memcheck, a memory error detector
==7345== Copyright (C) 2002-2011, and GNU GPL'd, by Julian Seward et al.
==7345== Using Valgrind-3.6.1 and LibVEX; rerun with -h for copyright info
==7345== Command: ls
==7345==
KandroidClient
KandroidServer
valgrind
==7345==
==7345== HEAP SUMMARY:
==7345==   in use at exit: 1,024 bytes in 1 blocks
==7345== total heap usage: 6 allocs, 5 frees, 5,295 bytes allocated
==7345==
==7345== LEAK SUMMARY:
==7345==   definitely lost: 0 bytes in 0 blocks
==7345==   indirectly lost: 0 bytes in 0 blocks
==7345==   possibly lost: 0 bytes in 0 blocks
==7345==   still reachable: 1,024 bytes in 1 blocks
==7345==             suppressed: 0 bytes in 0 blocks
==7345== Rerun with --leak-check=full to see details of leaked memory
==7345==
==7345== For counts of detected and suppressed errors, rerun with: -v
==7345== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
shell@android:/data/local/inst/bin $
```

# Valgrind Issues : No Swap and OOM Killer



# Memory Analysis Tools & Memory Management Methods

|         | General  | Android + External Tools   |
|---------|--|--|
| system  | <ul style="list-style-type: none"><li>• /proc/meminfo</li><li>• top</li></ul>  | <ul style="list-style-type: none"><li>• Low memory Killer</li><li>• DDMS SysInfo</li></ul>   |
| process | <ul style="list-style-type: none"><li>• /proc/&lt;pid&gt;/statm, maps, smaps</li><li>• Vss, Rss, Pss, Uss</li></ul>              | <ul style="list-style-type: none"><li>• procrank</li><li>• dumpsys meminfo</li><li>• smem</li></ul>  |
| Java    | Automatic Garbage Collection <ul style="list-style-type: none"><li>• Memory Leak Issues</li></ul>                                | <ul style="list-style-type: none"><li>• MAT</li><li>• Android GC - Native Bitmap Heap</li><li>• android.os.Debug.MemoryInfo</li><li>• ActivityManager.getProcessMemoryInfo</li></ul> |
| JNI     | Native to java object reference <ul style="list-style-type: none"><li>• Global Ref vs. Local Ref</li></ul>                       | <ul style="list-style-type: none"><li>• setprop dalvik.vm.checkjni true</li><li>• setprop dalvik.vm.jniopts forcetcopy</li><li>• setprop debug.checkjni 1</li></ul>                  |
| C++     | Smart Pointer <ul style="list-style-type: none"><li>• auto_ptr, scoped_ptr, unique_ptr,</li><li>• shared_ptr, weak_ptr</li></ul> | <ul style="list-style-type: none"><li>• valgrind</li><li>• sp, wp</li></ul>  |
| C       | glibc memory manager <ul style="list-style-type: none"><li>• ptmalloc (dlmalloc + pthread)</li></ul>                             | <ul style="list-style-type: none"><li>• valgrind</li><li>• bionic libc (dlmalloc)</li><li>• libc.debug.malloc</li></ul>  |

### AddressSanitizer vs. Valgrind (Memcheck)

|            | AddressSanitizer | Valgrind/Memcheck | Dr. Memory     | Mudflap       | Guard Page |
|------------|------------------|-------------------|----------------|---------------|------------|
| technology | CTI              | DBI               | DBI            | CTI           | Library    |
| ARCH       | x86              | x86,ARM,PPC       | x86            | all(?)        | all(?)     |
| OS         | Linux, Mac       | Linux, Mac        | Windows, Linux | Linux, Mac(?) | All (1)    |
| Slowdown   | 2x               | 20x               | 10x            | 2x-40x        | ?          |
| Detects:   |                  |                   |                |               |            |
| Heap OOB   | yes              | yes               | yes            | yes           | some       |
| Stack OOB  | yes              | no                | no             | some          | no         |
| Global OOB | yes              | no                | no             | ?             | no         |
| UAF        | yes              | yes               | yes            | yes           | yes        |
| UAR        | some             | no                | no             | no            | no         |
| UMR        | no               | yes               | yes            | ?             | no         |
| Leaks      | not yet          | yes               | yes            | ?             | no         |

**DBI:** dynamic binary instrumentation / **CTI:** compile-time instrumentation

**UMR:** uninitialized memory reads / **UAF:** use-after-free (aka dangling pointer)

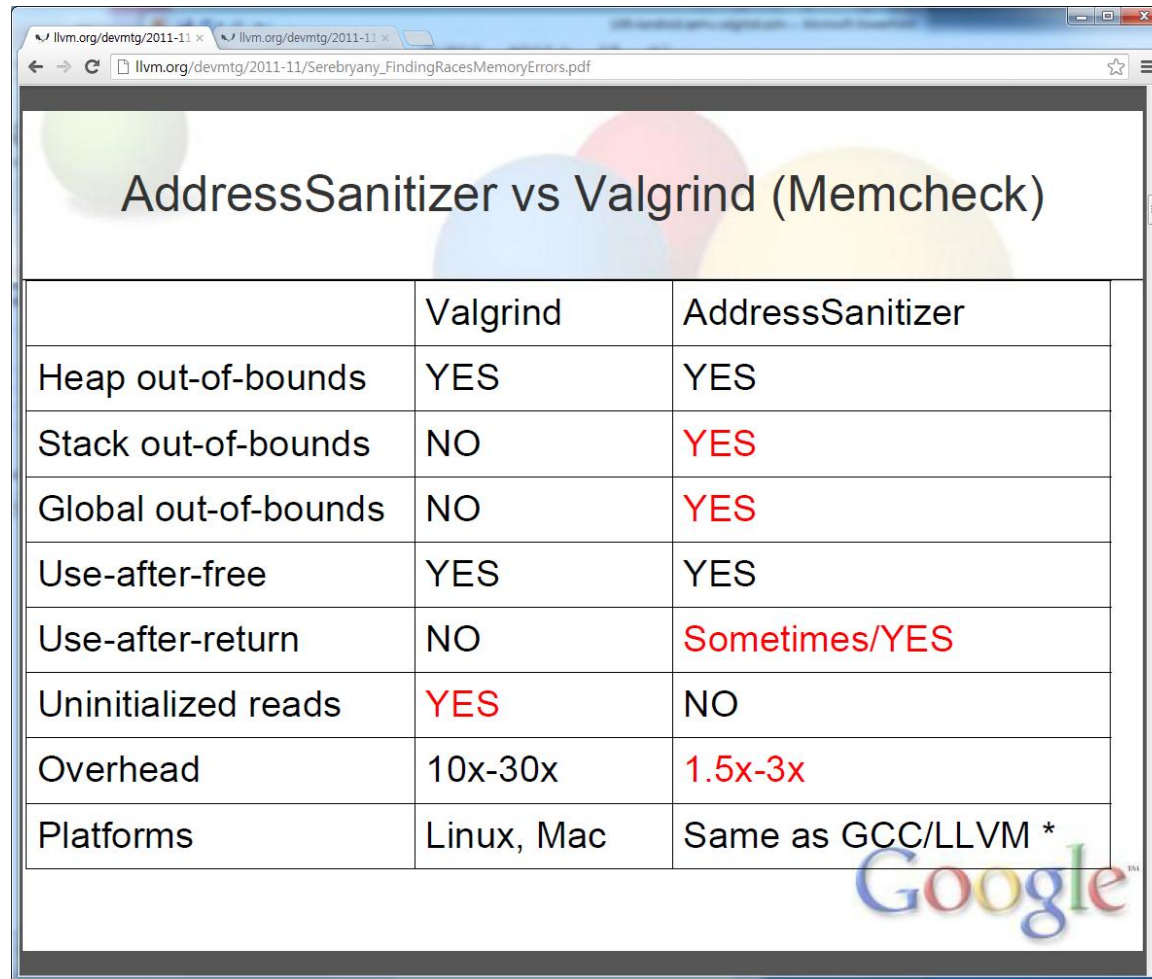
**UAR:** use-after-return / **OOB:** out-of-bounds

**x86:** includes 32- and 64-bit.

**Guard Page:** a family of memory error detectors (Electric fence or DUMA on Linux, Page Heap on Windows, Guard Malloc in Mac)

source : <http://code.google.com/p/address-sanitizer/wiki/ComparisonOfMemoryTools>

### AddressSanitizer vs. Valgrind (Memcheck)



The image is a screenshot of a web browser window displaying a PDF document. The browser's address bar shows the URL: [http://llvm.org/devmtg/2011-11/Serebryany\\_FindingRacesMemoryErrors.pdf](http://llvm.org/devmtg/2011-11/Serebryany_FindingRacesMemoryErrors.pdf). The PDF content features a title 'AddressSanitizer vs Valgrind (Memcheck)' at the top, followed by a comparison table. The table compares various memory error detection capabilities between Valgrind and AddressSanitizer. The background of the PDF has a decorative pattern of overlapping colored circles. A Google logo is visible in the bottom right corner of the PDF page.

|                      | Valgrind   | AddressSanitizer   |
|----------------------|------------|--------------------|
| Heap out-of-bounds   | YES        | YES                |
| Stack out-of-bounds  | NO         | YES                |
| Global out-of-bounds | NO         | YES                |
| Use-after-free       | YES        | YES                |
| Use-after-return     | NO         | Sometimes/YES      |
| Uninitialized reads  | YES        | NO                 |
| Overhead             | 10x-30x    | 1.5x-3x            |
| Platforms            | Linux, Mac | Same as GCC/LLVM * |

source : [http://llvm.org/devmtg/2011-11/Serebryany\\_FindingRacesMemoryErrors.pdf](http://llvm.org/devmtg/2011-11/Serebryany_FindingRacesMemoryErrors.pdf)

## Advanced Topics (1) – LLVM AddressSanitizer

---

```
$ svn co http://llvm.org/svn/llvm-project/llvm/trunk llvm
$ cd llvm/tools ; svn co http://llvm.org/svn/llvm-project/cfe/trunk clang
$ cd ../projects ; svn co http://llvm.org/svn/llvm-project/compiler-rt/trunk compiler-rt
$ svn co http://llvm.org/svn/llvm-project/test-suite/trunk test-suite
$ cd ../../ ; mkdir build ; cd build/ ; ../llvm/configure ; make
$ export PATH=<your_path>/Release+Asserts/bin/:$PATH:.
```

```
$ cd ../tools/clang
$ svn co http://src.chromium.org/svn/trunk/src/tools/clang/scripts
$ cd ../../
$ tools/clang/scripts/update.sh
```

```
$ cat > test/use-after-free.c
#include <stdlib.h>
int main() {
    char *x = (char*)malloc(10 * sizeof(char*));
    free(x);
    return x[5];
}

$ clang -faddress-sanitizer -O1 -fno-omit-frame-pointer -g tests/use-after-free.c
```

## Advanced Topics (1) – LLVM AddressSanitizer

---

```
$ ./a.out
```

```
=====
```

```
==8069== ERROR: AddressSanitizer:
```

```
heap-use-after-free on address 0x7f466f3ebf45 at pc 0x409284 bp 0x7fff3dd175f0 sp 0x7fff3dd175e8
```

```
READ of size 1 at 0x7f466f3ebf45 thread T0
```

```
  #0 0x409283 (/home/jsyang/llvm/llvm/a.out+0x409283)
```

```
  #1 0x7f466e85ec4c (/lib/libc-2.11.1.so+0x1ec4c)
```

```
0x7f466f3ebf45 is located 5 bytes inside of 80-byte region [0x7f466f3ebf40,0x7f466f3ebf90)
```

```
freed by thread T0 here:
```

```
  #0 0x4092d0 (/home/jsyang/llvm/llvm/a.out+0x4092d0)
```

```
  #1 0x40924a (/home/jsyang/llvm/llvm/a.out+0x40924a)
```

```
  #2 0x7f466e85ec4c (/lib/libc-2.11.1.so+0x1ec4c)
```

```
previously allocated by thread T0 here:
```

```
  #0 0x409390 (/home/jsyang/llvm/llvm/a.out+0x409390)
```

```
  #1 0x40923f (/home/jsyang/llvm/llvm/a.out+0x40923f)
```

```
  #2 0x7f466e85ec4c (/lib/libc-2.11.1.so+0x1ec4c)
```



### What Every Android Programmer Should Know About Memory?

- Paging : Contiguous Virtual Address vs. Noncontiguous Physical Address
- MMU : hardware vs. software
- Page : Clean/Dirty/Locked, File Mapping vs. Anonymous Mapping  
Reference Counted,
- Page Table and Protection Bit :  
PRESENT, PROTNONE, RW, USER, DIRTY, ACCESSED
- OOM Killer vs. Low Memory Killer
- Shared Memory vs. Shared Memory Virtual File System
- SYS V Shm, POSIX Shm, Berkeley Shm vs. Anonymous Shared Memory
- Process Address Segment : Code, Data, BSS, Heap, Stack, Mmap
- Memory Allocation : Static vs. Dynamic, Global vs. Stack
- malloc is a library function, not a system call
- Various malloc library implementation : dlmalloc, ptmalloc, ...
- Android C++ reference management :  
sp, wp – template + operator overloading
- Android Java GC includes the native Bitmap Heap.



Q & A



[www.kandroid.org](http://www.kandroid.org)